

# Mandrake Coordinator: Resource Management for Edge Clouds

Masters Presentation by Kyle Carson

Advisors: Chandra Krintz, Rich Wolski

Co-contributor: John Thomason

UCSB RACELab

03/12/19



# Background & Motivation

- The usage of Internet of Things (IoT) devices is rapidly increasing
- Deployments of such devices need computational infrastructure to support them
- Traditional “cloud computing” infrastructure can be unsuitable due to connectivity, latency, and data transfer constraints
- Need a new infrastructure paradigm, “edge clouds” offer a solution to these issues



# Edge Clouds

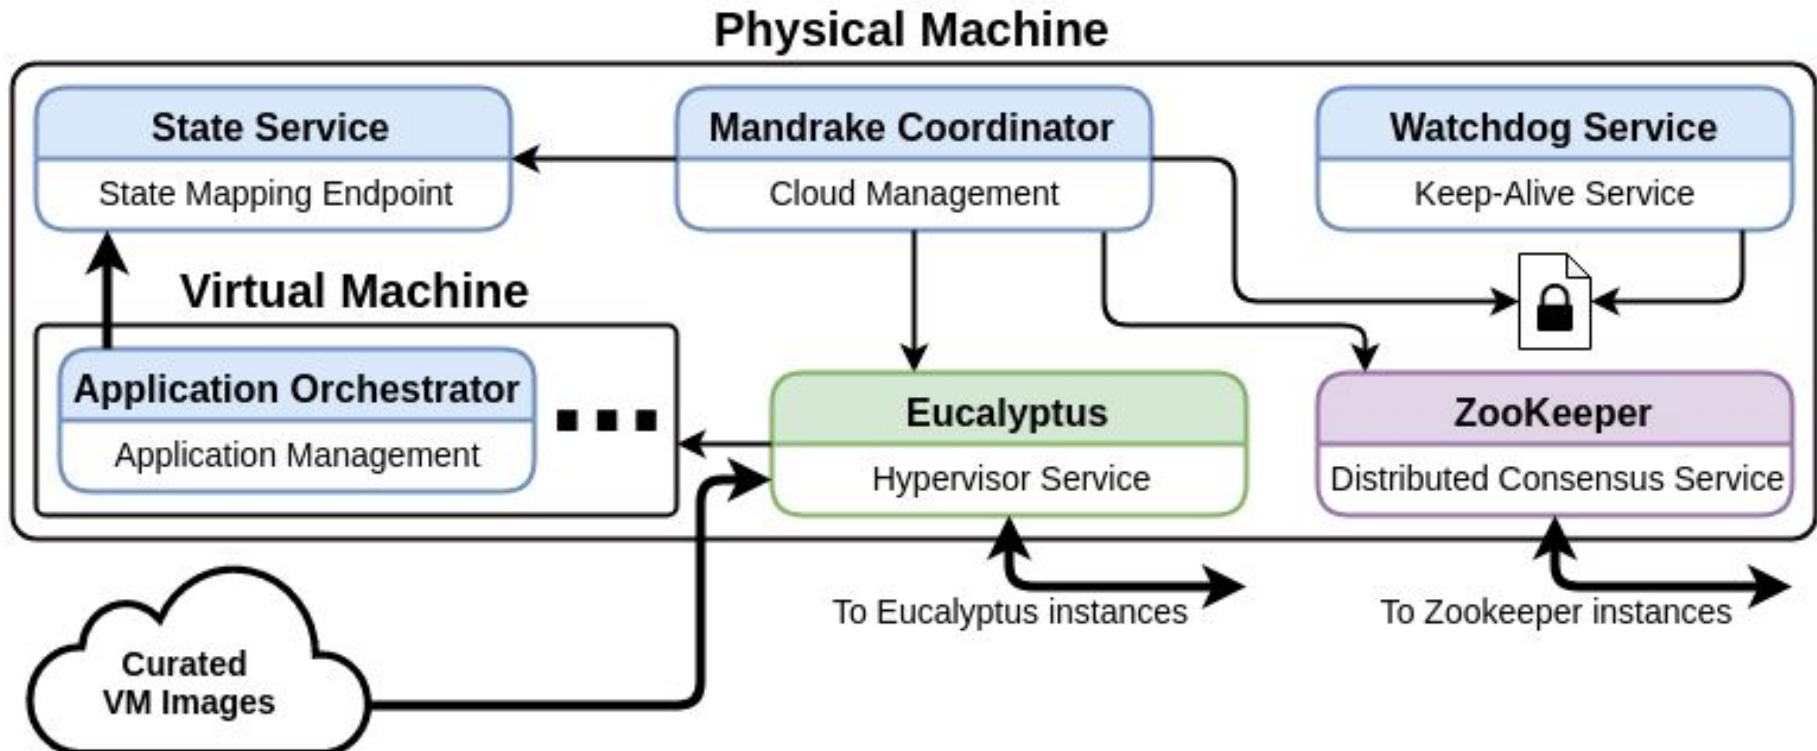
- Edge clouds are cloud computing infrastructure located at the “network edge”
- Enable real-time data processing with low latency response times by co-locating the infrastructure with the devices
- Additional issues including:
  - Hardware Failure
  - Network Partitioning
  - Delayed human maintenance



# Mandrake

- Mandrake is a software system which provides automated infrastructure & application management
- Reprovisions virtual machines (VMs) in response to failures and recoveries
- Informs VM-level software services of the VM-to-Host mapping, breaking the traditional cloud model of abstraction
- VM-level services manage user application, enabling it to make more informed decisions

# Architecture



# Mandrake Coordinator

- The Mandrake Coordinator (MC) discovers information about the hosts/VMs and keeps track of their mapping
- Exposes this map through a set of APIs
- A user-defined set of configuration controls the number of VMs, the way in which they scale, the images to use, etc.
- Controls the allocation/deallocation of VMs in response to failures and recoveries
- Makes guarantees about the liveness of the system, ensures resource allocation mechanisms will provide more reliability

# Experiment Setup


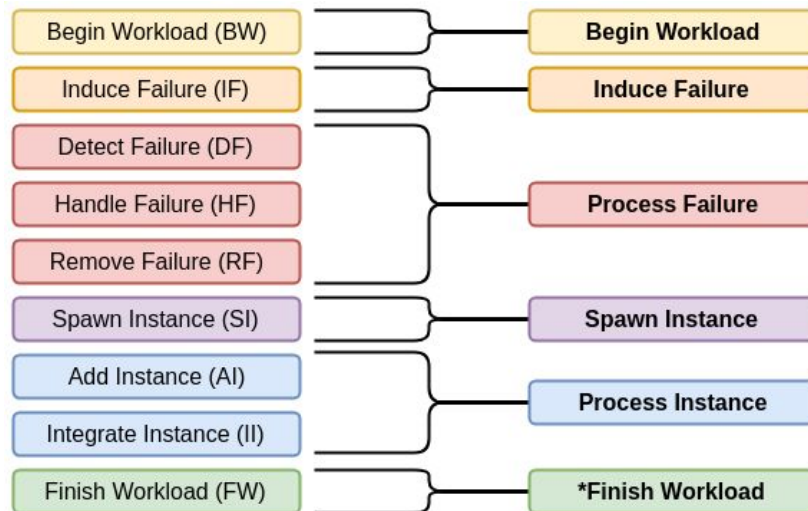
- 9 physical hosts with a target of 8 VMs, each with 2 CPU cores, 4GB of memory, and 100 GB of disk space
- Hadoop used as the case study user application
- HiBench used as the data processing benchmark
- Single failure simulated during each experiment run, 50 runs for each type of experiment





# Mandrake Coordinator Timeline

- Outline of key events during an experiment run
- The events can be clustered based on relative timings to represent broader “phases”






## Results: Performance

- Mandrake Coordinator took an average of...
  - ~15 seconds to detect failures
  - ~6 seconds to handle a failure/update the mapping
  - ~65 seconds for end-to-end reconfiguration of the infrastructure
  - ~95 seconds for end-to-end reconfiguration of the entire system (infrastructure + application ingestion)

| Phase           | Event | Avg. Time (seconds) |
|-----------------|-------|---------------------|
| Induce Failure  | IF    | 0.1                 |
| Process Failure | DF    | 14.6                |
|                 | HF    | 5.5                 |
|                 | RF    | 1.1                 |
| Spawn Instance  | SI    | 43.6                |

## Results: Failure Resilience

- Can sustain one or more failures every ~6 seconds and achieve a consistent mapping of the system at the infrastructure level
- Each failure “event” can be <50% of the machines, will reconfigure and continue to operate at degraded capacity



## Takeaways & Conclusions

- It is possible to significantly enhance the reliability and availability of a system such as an edge cloud without ruining the performance
- Can even improve long-term performance by guaranteeing functionality (at degraded capacity) when the cloud is in a damaged state
- Merit to breaking the abstraction of the cloud when dealing with resource constrained environments
- Overall, Mandrake allows user applications to make much more informed decisions and preserve in the face of failures

# Thank You!



Chandra Krintz



Rich Wolski



John Thomason



Markus Mock

...and everyone from RACELab, Port Hueneme Naval Base, the CS Dept. & UCSB

